

Predicting the sales price of bulldozers using Machine learning algorithm

1. Definition

How well can we predict the future sale price of a bulldozer, given its characteristics and previous examples of how much similar bulldozers have been sold for?

2. Data

The data we have taken from kaggle : <https://www.kaggle.com/c/bluebook-for-bulldozers/data>
(<https://www.kaggle.com/c/bluebook-for-bulldozers/data>).

There are 3 main datasets:

- Train.csv is the training set, which contains data through the end of 2011.
- Valid.csv is the validation set, which contains data from January 1, 2012 - April 30, 2012 You make predictions on this set throughout the majority of the competition. Your score on this set is used to create the public leaderboard.
- Test.csv is the test set, which won't be released until the last week of the competition. It contains data from May 1, 2012 - November 2012. Your score on the test set determines your final rank for the competition.

3. Evaluation

The evaluation metric for this competition is the RMSLE (root mean squared log error) between the actual and predicted auction prices.

For more on the evaluation of this project check : <https://www.kaggle.com/c/bluebook-for-bulldozers/overview/evaluation> (<https://www.kaggle.com/c/bluebook-for-bulldozers/overview/evaluation>)

Note. The goal for most regression evaluation metrics is to minimise the error. For example, our goal for this project is to build a machine learning model which minimises RMSLE.

4. Features

Kaggle provides a data dictionary detailing all the features of the dataset.

<https://www.kaggle.com/c/bluebook-for-bulldozers/data?select=Data+Dictionary.xlsx>
(<https://www.kaggle.com/c/bluebook-for-bulldozers/data?select=Data+Dictionary.xlsx>)

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import sklearn
```

```
In [120]: # Importing training and validation sets
df=pd.read_csv("data/TrainAndValid.csv", low_memory=False)
```

```
In [121]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 412698 entries, 0 to 412697
Data columns (total 53 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   SalesID                              412698 non-null  int64
1   SalePrice                            412698 non-null  float64
2   MachineID                            412698 non-null  int64
3   ModelID                              412698 non-null  int64
4   datasource                           412698 non-null  int64
5   auctioneerID                         392562 non-null  float64
6   YearMade                             412698 non-null  int64
7   MachineHoursCurrentMeter             147504 non-null  float64
8   UsageBand                            73670 non-null   object
9   saledate                             412698 non-null  object
10  fiModelDesc                           412698 non-null  object
11  fiBaseModel                           412698 non-null  object
12  fiSecondaryDesc                       271971 non-null  object
13  fiModelSeries                         58667 non-null   object
14  fiModelDescriptor                     74816 non-null   object
15  ProductSize                           196093 non-null  object
16  fiProductClassDesc                   412698 non-null  object
17  state                                 412698 non-null  object
18  ProductGroup                         412698 non-null  object
19  ProductGroupDesc                     412698 non-null  object
20  Drive_System                         107087 non-null  object
21  Enclosure                             412364 non-null  object
22  Forks                                 197715 non-null  object
23  Pad_Type                             81096 non-null   object
24  Ride_Control                         152728 non-null  object
25  Stick                                81096 non-null   object
26  Transmission                         188007 non-null  object
27  Turbocharged                         81096 non-null   object
28  Blade_Extension                      25983 non-null   object
29  Blade_Width                          25983 non-null   object
30  Enclosure_Type                       25983 non-null   object
31  Engine_Horsepower                   25983 non-null   object
32  Hydraulics                           330133 non-null  object
33  Pushblock                            25983 non-null   object
34  Ripper                               106945 non-null  object
35  Scarifier                            25994 non-null   object
36  Tip_Control                          25983 non-null   object
37  Tire_Size                            97638 non-null   object
38  Coupler                              220679 non-null  object
39  Coupler_System                       44974 non-null   object
40  Grouser_Tracks                       44875 non-null   object
41  Hydraulics_Flow                      44875 non-null   object
42  Track_Type                           102193 non-null  object
43  Undercarriage_Pad_Width              102916 non-null  object
44  Stick_Length                         102261 non-null  object
45  Thumb                                102332 non-null  object
46  Pattern_Changer                      102261 non-null  object
47  Grouser_Type                         102193 non-null  object
48  Backhoe_Mounting                     80712 non-null   object
49  Blade_Type                           81875 non-null   object
```

50	Travel_Controls	81877 non-null	object
51	Differential_Type	71564 non-null	object
52	Steering_Controls	71522 non-null	object

dtypes: float64(3), int64(5), object(45)
memory usage: 166.9+ MB

```
In [122]: df.isna().sum()
```

```
Out[122]: SalesID                0
SalePrice                0
MachineID                0
ModelID                  0
datasource                0
auctioneerID             20136
YearMade                  0
MachineHoursCurrentMeter  265194
UsageBand                 339028
saledate                  0
fiModelDesc               0
fiBaseModel               0
fiSecondaryDesc           140727
fiModelSeries             354031
fiModelDescriptor         337882
ProductSize               216605
fiProductClassDesc        0
state                     0
ProductGroup              0
ProductGroupDesc          0
Drive_System              305611
Enclosure                  334
Forks                     214983
Pad_Type                   331602
Ride_Control               259970
Stick                     331602
Transmission              224691
Turbocharged              331602
Blade_Extension           386715
Blade_Width               386715
Enclosure_Type            386715
Engine_Horsepower         386715
Hydraulics                 82565
Pushblock                 386715
Ripper                    305753
Scarifier                 386704
Tip_Control               386715
Tire_Size                 315060
Coupler                   192019
Coupler_System            367724
Grouser_Tracks            367823
Hydraulics_Flow           367823
Track_Type                310505
Undercarriage_Pad_Width   309782
Stick_Length              310437
Thumb                     310366
Pattern_Changer           310437
Grouser_Type              310505
Backhoe_Mounting          331986
Blade_Type                330823
Travel_Controls           330821
Differential_Type         341134
Steering_Controls         341176
dtype: int64
```

```
In [123]: df.saledate
```


```
Out[123]: 0          11/16/2006 0:00
1          3/26/2004 0:00
2          2/26/2004 0:00
3          5/19/2011 0:00
4          7/23/2009 0:00
...
412693     3/7/2012 0:00
412694     1/28/2012 0:00
412695     1/28/2012 0:00
412696     3/7/2012 0:00
412697     1/28/2012 0:00
Name: saledate, Length: 412698, dtype: object
```

```
In [124]: df.saledate.dtype
```

```
Out[124]: dtype('O')
```

```
In [125]: df.SalePrice.plot.hist()
```

```
Out[125]: <matplotlib.axes._subplots.AxesSubplot at 0x1816ebcf1c8>
```



Parsing dates

When we work with time series data, we want to enrich the time & date component as much as possible.

We can do that by telling which of our columns has dates in it using the `parse_dates` parameter.

```
In [126]: # Import data again but this time parse dates
df=pd.read_csv("data/TrainAndValid.csv",
               low_memory=False,
               parse_dates=["saledate"])
```

```
In [127]: df.saledate.dtype
```


```
Out[127]: dtype('<M8[ns]')
```

```
In [128]: df.saledate[:1000]
```

```
Out[128]: 0      2006-11-16
1      2004-03-26
2      2004-02-26
3      2011-05-19
4      2009-07-23
...
995    2009-07-16
996    2007-06-14
997    2005-09-22
998    2005-07-28
999    2011-06-16
Name: saledate, Length: 1000, dtype: datetime64[ns]
```

```
In [129]: fig, ax = plt.subplots()
ax.scatter(df["saledate"][:1000], df["SalePrice"][:1000])
```

```
Out[129]: <matplotlib.collections.PathCollection at 0x18137af4308>
```




In [130]: `df.head()`

Out[130]:

	SalesID	SalePrice	MachineID	ModelID	datasource	auctioneerID	YearMade	MachineHoursCur
0	1139246	66000.0	999089	3157	121	3.0	2004	
1	1139248	57000.0	117657	77	121	3.0	1996	
2	1139249	10000.0	434808	7009	121	3.0	2001	
3	1139251	38500.0	1026470	332	121	3.0	2001	
4	1139253	11000.0	1057373	17311	121	3.0	2007	

5 rows × 53 columns



In [131]: df.head().T

Out[131]:

	0	1	2	3	4
SalesID	1139246	1139248	1139249	1139251	1139253
SalePrice	66000	57000	10000	38500	11000
MachineID	999089	117657	434808	1026470	1057373
ModelID	3157	77	7009	332	17311
datasource	121	121	121	121	121
auctioneerID	3	3	3	3	3
YearMade	2004	1996	2001	2001	2007
MachineHoursCurrentMeter	68	4640	2838	3486	722
UsageBand	Low	Low	High	High	Medium
saledate	2006-11-16 00:00:00	2004-03-26 00:00:00	2004-02-26 00:00:00	2011-05-19 00:00:00	2009-07-23 00:00:00
fiModelDesc	521D	950FII	226	PC120-6E	S175
fiBaseModel	521	950	226	PC120	S175
fiSecondaryDesc	D	F	NaN	NaN	NaN
fiModelSeries	NaN	II	NaN	-6E	NaN
fiModelDescriptor	NaN	NaN	NaN	NaN	NaN
ProductSize	NaN	Medium	NaN	Small	NaN
fiProductClassDesc	Wheel Loader - 110.0 to 120.0 Horsepower	Wheel Loader - 150.0 to 175.0 Horsepower	Skid Steer Loader - 1351.0 to 1601.0 Lb Operat...	Hydraulic Excavator, Track - 12.0 to 14.0 Metr...	Skid Steer Loader - 1601.0 to 1751.0 Lb Operat...
state	Alabama	North Carolina	New York	Texas	New York
ProductGroup	WL	WL	SSL	TEX	SSL
ProductGroupDesc	Wheel Loader	Wheel Loader	Skid Steer Loaders	Track Excavators	Skid Steer Loaders
Drive_System	NaN	NaN	NaN	NaN	NaN
Enclosure	EROPS w AC	EROPS w AC	OROPS	EROPS w AC	EROPS
Forks	None or Unspecified	None or Unspecified	None or Unspecified	NaN	None or Unspecified
Pad_Type	NaN	NaN	NaN	NaN	NaN
Ride_Control	None or Unspecified	None or Unspecified	NaN	NaN	NaN
Stick	NaN	NaN	NaN	NaN	NaN
Transmission	NaN	NaN	NaN	NaN	NaN
Turbocharged	NaN	NaN	NaN	NaN	NaN
Blade_Extension	NaN	NaN	NaN	NaN	NaN

	0	1	2	3	4
Blade_Width	NaN	NaN	NaN	NaN	NaN
Enclosure_Type	NaN	NaN	NaN	NaN	NaN
Engine_Horsepower	NaN	NaN	NaN	NaN	NaN
Hydraulics	2 Valve	2 Valve	Auxiliary	2 Valve	Auxiliary
Pushblock	NaN	NaN	NaN	NaN	NaN
Ripper	NaN	NaN	NaN	NaN	NaN
Scarifier	NaN	NaN	NaN	NaN	NaN
Tip_Control	NaN	NaN	NaN	NaN	NaN
Tire_Size	None or Unspecified	23.5	NaN	NaN	NaN
Coupler	None or Unspecified	None or Unspecified	None or Unspecified	None or Unspecified	None or Unspecified
Coupler_System	NaN	NaN	None or Unspecified	NaN	None or Unspecified
Grouser_Tracks	NaN	NaN	None or Unspecified	NaN	None or Unspecified
Hydraulics_Flow	NaN	NaN	Standard	NaN	Standard
Track_Type	NaN	NaN	NaN	NaN	NaN
Undercarriage_Pad_Width	NaN	NaN	NaN	NaN	NaN
Stick_Length	NaN	NaN	NaN	NaN	NaN
Thumb	NaN	NaN	NaN	NaN	NaN
Pattern_Changer	NaN	NaN	NaN	NaN	NaN
Grouser_Type	NaN	NaN	NaN	NaN	NaN
Backhoe_Mounting	NaN	NaN	NaN	NaN	NaN
Blade_Type	NaN	NaN	NaN	NaN	NaN
Travel_Controls	NaN	NaN	NaN	NaN	NaN
Differential_Type	Standard	Standard	NaN	NaN	NaN
Steering_Controls	Conventional	Conventional	NaN	NaN	NaN

```
In [132]: df.saledate[:20]
```

```
Out[132]: 0      2006-11-16
1      2004-03-26
2      2004-02-26
3      2011-05-19
4      2009-07-23
5      2008-12-18
6      2004-08-26
7      2005-11-17
8      2009-08-27
9      2007-08-09
10     2008-08-21
11     2006-08-24
12     2005-10-20
13     2006-01-26
14     2006-01-03
15     2006-11-16
16     2007-06-14
17     2010-01-28
18     2006-03-09
19     2005-11-17
Name: saledate, dtype: datetime64[ns]
```

Sort DataFrame by saledate

When working with time series data, its a good practice to sort the data in ascending order of date.

```
In [133]: df.sort_values(by=["saledate"],inplace=True,ascending=True)
df.saledate.head(20)
```

```
Out[133]: 205615    1989-01-17
274835    1989-01-31
141296    1989-01-31
212552    1989-01-31
62755     1989-01-31
54653     1989-01-31
81383     1989-01-31
204924    1989-01-31
135376    1989-01-31
113390    1989-01-31
113394    1989-01-31
116419    1989-01-31
32138     1989-01-31
127610    1989-01-31
76171     1989-01-31
127000    1989-01-31
128130    1989-01-31
127626    1989-01-31
55455     1989-01-31
55454     1989-01-31
Name: saledate, dtype: datetime64[ns]
```

In [134]: df.head().T

Out[134]:

	205615	274835	141296	212552	62755
SalesID	1646770	1821514	1505138	1671174	1329056
SalePrice	9500	14000	50000	16000	22000
MachineID	1126363	1194089	1473654	1327630	1336053
ModelID	8434	10150	4139	8591	4089
datasource	132	132	132	132	132
auctioneerID	18	99	99	99	99
YearMade	1974	1980	1978	1980	1984
MachineHoursCurrentMeter	NaN	NaN	NaN	NaN	NaN
UsageBand	NaN	NaN	NaN	NaN	NaN
saledate	1989-01-17 00:00:00	1989-01-31 00:00:00	1989-01-31 00:00:00	1989-01-31 00:00:00	1989-01-31 00:00:00
fiModelDesc	TD20	A66	D7G	A62	D3B
fiBaseModel	TD20	A66	D7	A62	D3
fiSecondaryDesc	NaN	NaN	G	NaN	B
fiModelSeries	NaN	NaN	NaN	NaN	NaN
fiModelDescriptor	NaN	NaN	NaN	NaN	NaN
ProductSize	Medium	NaN	Large	NaN	NaN
fiProductClassDesc	Track Type Tractor, Dozer - 105.0 to 130.0 Hor...	Wheel Loader - 120.0 to 135.0 Horsepower	Track Type Tractor, Dozer - 190.0 to 260.0 Hor...	Wheel Loader - Unidentified	Track Type Tractor, Dozer - 20.0 to 75.0 Horse...
state	Texas	Florida	Florida	Florida	Florida
ProductGroup	TTT	WL	TTT	WL	TTT
ProductGroupDesc	Track Type Tractors	Wheel Loader	Track Type Tractors	Wheel Loader	Track Type Tractors
Drive_System	NaN	NaN	NaN	NaN	NaN
Enclosure	OROPS	OROPS	OROPS	EROPS	OROPS
Forks	NaN	None or Unspecified	NaN	None or Unspecified	NaN
Pad_Type	NaN	NaN	NaN	NaN	NaN
Ride_Control	NaN	None or Unspecified	NaN	None or Unspecified	NaN
Stick	NaN	NaN	NaN	NaN	NaN
Transmission	Direct Drive	NaN	Standard	NaN	Standard
Turbocharged	NaN	NaN	NaN	NaN	NaN
Blade_Extension	NaN	NaN	NaN	NaN	NaN
Blade_Width	NaN	NaN	NaN	NaN	NaN

	205615	274835	141296	212552	62755
Enclosure_Type	NaN	NaN	NaN	NaN	NaN
Engine_Horsepower	NaN	NaN	NaN	NaN	NaN
Hydraulics	2 Valve	2 Valve	2 Valve	2 Valve	2 Valve
Pushblock	NaN	NaN	NaN	NaN	NaN
Ripper	None or Unspecified	NaN	None or Unspecified	NaN	None or Unspecified
Scarifier	NaN	NaN	NaN	NaN	NaN
Tip_Control	NaN	NaN	NaN	NaN	NaN
Tire_Size	NaN	None or Unspecified	NaN	None or Unspecified	NaN
Coupler	NaN	None or Unspecified	NaN	None or Unspecified	NaN
Coupler_System	NaN	NaN	NaN	NaN	NaN
Grouser_Tracks	NaN	NaN	NaN	NaN	NaN
Hydraulics_Flow	NaN	NaN	NaN	NaN	NaN
Track_Type	NaN	NaN	NaN	NaN	NaN
Undercarriage_Pad_Width	NaN	NaN	NaN	NaN	NaN
Stick_Length	NaN	NaN	NaN	NaN	NaN
Thumb	NaN	NaN	NaN	NaN	NaN
Pattern_Changer	NaN	NaN	NaN	NaN	NaN
Grouser_Type	NaN	NaN	NaN	NaN	NaN
Backhoe_Mounting	None or Unspecified	NaN	None or Unspecified	NaN	None or Unspecified
Blade_Type	Straight	NaN	Straight	NaN	PAT
Travel_Controls	None or Unspecified	NaN	None or Unspecified	NaN	Lever
Differential_Type	NaN	Standard	NaN	Standard	NaN
Steering_Controls	NaN	Conventional	NaN	Conventional	NaN


```
In [135]: ## Make a copy of original dataframe, so that when we manipulate the copy we still have the original
df_tmp=df.copy()
```

```
In [136]: df_tmp.head()
```

```
Out[136]:
```

	SalesID	SalePrice	MachineID	ModelID	datasource	auctioneerID	YearMade	MachineHou
205615	1646770	9500.0	1126363	8434	132	18.0	1974	
274835	1821514	14000.0	1194089	10150	132	99.0	1980	
141296	1505138	50000.0	1473654	4139	132	99.0	1978	
212552	1671174	16000.0	1327630	8591	132	99.0	1980	
62755	1329056	22000.0	1336053	4089	132	99.0	1984	

5 rows × 53 columns



```
In [137]: ## We have done some feature engineering which required in timeSeries kind of proc  
df_tmp["saleYear"] = df_tmp.saledate.dt.year  
df_tmp["saleMonth"] = df_tmp.saledate.dt.month  
df_tmp["saleDay"] = df_tmp.saledate.dt.day  
df_tmp["saleDayOfWeek"] = df_tmp.saledate.dt.dayofweek  
df_tmp["saleDayOfYear"] = df_tmp.saledate.dt.dayofyear
```

In [138]: df_tmp.head().T

Out[138]:

	205615	274835	141296	212552	62755
SalesID	1646770	1821514	1505138	1671174	1329056
SalePrice	9500	14000	50000	16000	22000
MachineID	1126363	1194089	1473654	1327630	1336053
ModelID	8434	10150	4139	8591	4089
datasource	132	132	132	132	132
auctioneerID	18	99	99	99	99
YearMade	1974	1980	1978	1980	1984
MachineHoursCurrentMeter	NaN	NaN	NaN	NaN	NaN
UsageBand	NaN	NaN	NaN	NaN	NaN
saledate	1989-01-17 00:00:00	1989-01-31 00:00:00	1989-01-31 00:00:00	1989-01-31 00:00:00	1989-01-31 00:00:00
fiModelDesc	TD20	A66	D7G	A62	D3B
fiBaseModel	TD20	A66	D7	A62	D3
fiSecondaryDesc	NaN	NaN	G	NaN	B
fiModelSeries	NaN	NaN	NaN	NaN	NaN
fiModelDescriptor	NaN	NaN	NaN	NaN	NaN
ProductSize	Medium	NaN	Large	NaN	NaN
fiProductClassDesc	Track Type Tractor, Dozer - 105.0 to 130.0 Hor...	Wheel Loader - 120.0 to 135.0 Horsepower	Track Type Tractor, Dozer - 190.0 to 260.0 Hor...	Wheel Loader - Unidentified	Track Type Tractor, Dozer - 20.0 to 75.0 Horse...
state	Texas	Florida	Florida	Florida	Florida
ProductGroup	TTT	WL	TTT	WL	TTT
ProductGroupDesc	Track Type Tractors	Wheel Loader	Track Type Tractors	Wheel Loader	Track Type Tractors
Drive_System	NaN	NaN	NaN	NaN	NaN
Enclosure	OROPS	OROPS	OROPS	EROPS	OROPS
Forks	NaN	None or Unspecified	NaN	None or Unspecified	NaN
Pad_Type	NaN	NaN	NaN	NaN	NaN
Ride_Control	NaN	None or Unspecified	NaN	None or Unspecified	NaN
Stick	NaN	NaN	NaN	NaN	NaN
Transmission	Direct Drive	NaN	Standard	NaN	Standard
Turbocharged	NaN	NaN	NaN	NaN	NaN
Blade_Extension	NaN	NaN	NaN	NaN	NaN
Blade_Width	NaN	NaN	NaN	NaN	NaN

	205615	274835	141296	212552	62755
Enclosure_Type	NaN	NaN	NaN	NaN	NaN
Engine_Horsepower	NaN	NaN	NaN	NaN	NaN
Hydraulics	2 Valve	2 Valve	2 Valve	2 Valve	2 Valve
Pushblock	NaN	NaN	NaN	NaN	NaN
Ripper	None or Unspecified	NaN	None or Unspecified	NaN	None or Unspecified
Scarifier	NaN	NaN	NaN	NaN	NaN
Tip_Control	NaN	NaN	NaN	NaN	NaN
Tire_Size	NaN	None or Unspecified	NaN	None or Unspecified	NaN
Coupler	NaN	None or Unspecified	NaN	None or Unspecified	NaN
Coupler_System	NaN	NaN	NaN	NaN	NaN
Grouser_Tracks	NaN	NaN	NaN	NaN	NaN
Hydraulics_Flow	NaN	NaN	NaN	NaN	NaN
Track_Type	NaN	NaN	NaN	NaN	NaN
Undercarriage_Pad_Width	NaN	NaN	NaN	NaN	NaN
Stick_Length	NaN	NaN	NaN	NaN	NaN
Thumb	NaN	NaN	NaN	NaN	NaN
Pattern_Changer	NaN	NaN	NaN	NaN	NaN
Grouser_Type	NaN	NaN	NaN	NaN	NaN
Backhoe_Mounting	None or Unspecified	NaN	None or Unspecified	NaN	None or Unspecified
Blade_Type	Straight	NaN	Straight	NaN	PAT
Travel_Controls	None or Unspecified	NaN	None or Unspecified	NaN	Lever
Differential_Type	NaN	Standard	NaN	Standard	NaN
Steering_Controls	NaN	Conventional	NaN	Conventional	NaN
saleYear	1989	1989	1989	1989	1989
saleMonth	1	1	1	1	1
saleDay	17	31	31	31	31
saleDayOfWeek	1	1	1	1	1
saleDayOfYear	17	31	31	31	31

```
In [139]: ## Now we have saleDay, saleMonth, saleYear, So we don't require saledate anymore
df_tmp.drop(labels="saledate",axis=1,inplace=True)
```



```
In [140]: df_tmp.head(20).T
```

```
Out[140]:
```

	205615	274835	141296	212552	62755	5
SalesID	1646770	1821514	1505138	1671174	1329056	130
SalePrice	9500	14000	50000	16000	22000	2
MachineID	1126363	1194089	1473654	1327630	1336053	118
ModelID	8434	10150	4139	8591	4089	
datasource	132	132	132	132	132	
auctioneerID	18	99	99	99	99	
YearMade	1974	1980	1978	1980	1984	
MachineHoursCurrentMeter	NaN	NaN	NaN	NaN	NaN	
UsageBand	NaN	NaN	NaN	NaN	NaN	
fiModelDesc	TD20	A66	D7G	A62	D3B	
fiBaseModel	TD20	A66	D7	A62	D3	

```
In [141]: df_tmp.saledate
```

```
-----
AttributeError                                Traceback (most recent call last)
<ipython-input-141-0b8254724953> in <module>
----> 1 df_tmp.saledate

~\AppData\Roaming\Python\Python37\site-packages\pandas\core\generic.py in __getattribute__(self, name)
    5139         if self._info_axis._can_hold_identifiers_and_holds_name(name):
    5140             return self[name]
-> 5141         return object.__getattribute__(self, name)
    5142
    5143     def __setattr__(self, name: str, value) -> None:

AttributeError: 'DataFrame' object has no attribute 'saledate'
```

```
In [144]: df_tmp.state.value_counts()
```

```
Out[144]: Florida          67320
Texas          53110
California     29761
Washington     16222
Georgia        14633
Maryland       13322
Mississippi    13240
Ohio           12369
Illinois       11540
Colorado       11529
New Jersey     11156
North Carolina 10636
Tennessee     10298
Alabama        10292
Pennsylvania   10234
South Carolina  9951
Arizona        9364
New York       8639
Connecticut    8276
Minnesota      7885
Missouri       7178
Nevada         6932
Louisiana      6627
Kentucky       5351
Maine          5096
Indiana        4124
Arkansas       3933
New Mexico     3631
Utah           3046
Unspecified    2801
Wisconsin      2745
New Hampshire  2738
Virginia       2353
Idaho          2025
Oregon         1911
Michigan       1831
Wyoming        1672
Montana        1336
Iowa           1336
Oklahoma       1326
Nebraska       866
West Virginia  840
Kansas         667
Delaware       510
North Dakota   480
Alaska         430
Massachusetts  347
Vermont        300
South Dakota   244
Hawaii         118
Rhode Island   83
Puerto Rico   42
Washington DC  2
Name: state, dtype: int64
```

```
In [145]: len(df_tmp)
```

```
Out[145]: 412698
```

5. Modelling

We have done enough EDA (we could always do more) but let's start to do some model-driven EDA

```
In [146]: ## Let's build a machine learning model
from sklearn.ensemble import RandomForestRegressor
model=RandomForestRegressor(n_jobs=1,
                             random_state=42) # same as np.seed(42)
model.fit(df_tmp.drop("SalePrice", axis=1), df_tmp["SalePrice"])
```

```
-----
ValueError                                Traceback (most recent call last)
<ipython-input-146-b3ff0195a76f> in <module>
      3 model=RandomForestRegressor(n_jobs=1,
      4                             random_state=42) # same as np.seed(42)
----> 5 model.fit(df_tmp.drop("SalePrice", axis=1), df_tmp["SalePrice"])

~\AppData\Roaming\Python\Python37\site-packages\sklearn\ensemble\_forest.py in
fit(self, X, y, sample_weight)
    303         )
    304         X, y = self._validate_data(X, y, multi_output=True,
--> 305                                     accept_sparse="csc", dtype=DTYPE)
    306         if sample_weight is not None:
    307             sample_weight = _check_sample_weight(sample_weight, X)

~\AppData\Roaming\Python\Python37\site-packages\sklearn\base.py in _validate_da
ta(self, X, y, reset, validate_separately, **check_params)
    431         y = check_array(y, **check_y_params)
    432         else:
--> 433             X, y = check_X_y(X, y, **check_params)
    434             out = X, y
    435

~\AppData\Roaming\Python\Python37\site-packages\sklearn\utils\validation.py in
inner_f(*args, **kwargs)
    61         extra_args = len(args) - len(all_args)
    62         if extra_args <= 0:
---> 63             return f(*args, **kwargs)
    64
    65             # extra_args > 0

~\AppData\Roaming\Python\Python37\site-packages\sklearn\utils\validation.py in
check_X_y(X, y, accept_sparse, accept_large_sparse, dtype, order, copy, force_a
ll_finite, ensure_2d, allow_nd, multi_output, ensure_min_samples, ensure_min_fe
atures, y_numeric, estimator)
    876         ensure_min_samples=ensure_min_samples,
    877         ensure_min_features=ensure_min_features,
--> 878         estimator=estimator)
    879     if multi_output:
    880         y = check_array(y, accept_sparse='csr', force_all_finite=True,

~\AppData\Roaming\Python\Python37\site-packages\sklearn\utils\validation.py in
inner_f(*args, **kwargs)
    61         extra_args = len(args) - len(all_args)
    62         if extra_args <= 0:
---> 63             return f(*args, **kwargs)
    64
    65             # extra_args > 0

~\AppData\Roaming\Python\Python37\site-packages\sklearn\utils\validation.py in
check_array(array, accept_sparse, accept_large_sparse, dtype, order, copy, forc
```

```

e_all_finite, ensure_2d, allow_nd, ensure_min_samples, ensure_min_features, estimator)
671         array = array.astype(dtype, casting="unsafe", copy=False)
672     else:
--> 673         array = np.asarray(array, order=order, dtype=dtype)
674     except ComplexWarning as complex_warning:
675         raise ValueError("Complex data not supported\n"

```

```

~\AppData\Roaming\Python\Python37\site-packages\numpy\core\_asarray.py in asarray(a, dtype, order)
81
82     """
---> 83     return array(a, dtype, copy=False, order=order)
84
85

```

```

~\AppData\Roaming\Python\Python37\site-packages\pandas\core\generic.py in __array__(self, dtype)
1779
1780     def __array__(self, dtype=None) -> np.ndarray:
-> 1781         return np.asarray(self._values, dtype=dtype)
1782
1783     def __array_wrap__(self, result, context=None):

```

```

~\AppData\Roaming\Python\Python37\site-packages\numpy\core\_asarray.py in asarray(a, dtype, order)
81
82     """
---> 83     return array(a, dtype, copy=False, order=order)
84
85

```

ValueError: could not convert string to float: 'Low'

```
In [147]: df_tmp.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 412698 entries, 205615 to 409203
Data columns (total 57 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   SalesID                               412698 non-null  int64
1   SalePrice                             412698 non-null  float64
2   MachineID                             412698 non-null  int64
3   ModelID                               412698 non-null  int64
4   datasource                            412698 non-null  int64
5   auctioneerID                          392562 non-null  float64
6   YearMade                              412698 non-null  int64
7   MachineHoursCurrentMeter              147504 non-null  float64
8   UsageBand                             73670 non-null   object
9   fiModelDesc                           412698 non-null  object
10  fiBaseModel                           412698 non-null  object
11  fiSecondaryDesc                        271971 non-null  object
12  fiModelSeries                          58667 non-null   object
13  fiModelDescriptor                      74816 non-null   object
14  ProductSize                           196093 non-null  object
15  fiProductClassDesc                    412698 non-null  object
16  state                                 412698 non-null  object
17  ProductGroup                          412698 non-null  object
18  ProductGroupDesc                      412698 non-null  object
19  Drive_System                          107087 non-null  object
20  Enclosure                             412364 non-null  object
21  Forks                                 197715 non-null  object
22  Pad_Type                              81096 non-null   object
23  Ride_Control                          152728 non-null  object
24  Stick                                 81096 non-null   object
25  Transmission                          188007 non-null  object
26  Turbocharged                          81096 non-null   object
27  Blade_Extension                       25983 non-null   object
28  Blade_Width                           25983 non-null   object
29  Enclosure_Type                        25983 non-null   object
30  Engine_Horsepower                     25983 non-null   object
31  Hydraulics                            330133 non-null  object
32  Pushblock                             25983 non-null   object
33  Ripper                                106945 non-null  object
34  Scarifier                             25994 non-null   object
35  Tip_Control                           25983 non-null   object
36  Tire_Size                             97638 non-null   object
37  Coupler                               220679 non-null  object
38  Coupler_System                        44974 non-null   object
39  Grouser_Tracks                        44875 non-null   object
40  Hydraulics_Flow                       44875 non-null   object
41  Track_Type                            102193 non-null  object
42  Undercarriage_Pad_Width               102916 non-null  object
43  Stick_Length                          102261 non-null  object
44  Thumb                                 102332 non-null  object
45  Pattern_Changer                       102261 non-null  object
46  Grouser_Type                          102193 non-null  object
47  Backhoe_Mounting                      80712 non-null   object
48  Blade_Type                            81875 non-null   object
49  Travel_Controls                       81877 non-null   object
```

```
50 Differential_Type      71564 non-null  object
51 Steering_Controls     71522 non-null  object
52 saleYear              412698 non-null  int64
53 saleMonth             412698 non-null  int64
54 saleDay               412698 non-null  int64
55 saleDayOfWeek         412698 non-null  int64
56 saleDayOfYear         412698 non-null  int64
dtypes: float64(3), int64(10), object(44)
memory usage: 182.6+ MB
```

```
In [148]: df_tmp["UsageBand"].dtype
```

```
Out[148]: dtype('O')
```

```
In [149]: df_tmp.isna().sum()
```

```
Out[149]: SalesID                0
SalePrice                0
MachineID                0
ModelID                  0
datasource                0
auctioneerID             20136
YearMade                  0
MachineHoursCurrentMeter 265194
UsageBand                 339028
fiModelDesc               0
fiBaseModel               0
fiSecondaryDesc           140727
fiModelSeries             354031
fiModelDescriptor         337882
ProductSize               216605
fiProductClassDesc        0
state                     0
ProductGroup              0
ProductGroupDesc          0
Drive_System              305611
Enclosure                 334
Forks                     214983
Pad_Type                  331602
Ride_Control              259970
Stick                     331602
Transmission              224691
Turbocharged              331602
Blade_Extension           386715
Blade_Width               386715
Enclosure_Type            386715
Engine_Horsepower         386715
Hydraulics                82565
Pushblock                 386715
Ripper                    305753
Scarifier                 386704
Tip_Control               386715
Tire_Size                 315060
Coupler                   192019
Coupler_System            367724
Grouser_Tracks            367823
Hydraulics_Flow           367823
Track_Type                310505
Undercarriage_Pad_Width   309782
Stick_Length              310437
Thumb                     310366
Pattern_Changer           310437
Grouser_Type              310505
Backhoe_Mounting          331986
Blade_Type                330823
Travel_Controls           330821
Differential_Type         341134
Steering_Controls         341176
saleYear                  0
saleMonth                 0
saleDay                   0
```



```
saleDayOfWeek      0
saleDayOfYear       0
dtype: int64
```

Convert String to Categorical data

One way we can turn all our data into numbers is by converting them into pandas categories.

```
In [150]: pd.api.types.is_string_dtype(df_tmp["UsageBand"])
```

```
Out[150]: True
```

In [151]: *## Getting all the column names which has dtype as string*

```
for label,content in df_tmp.items():  
    if(pd.api.types.is_string_dtype(content)):  
        print(label)
```

UsageBand
fiModelDesc
fiBaseModel
fiSecondaryDesc
fiModelSeries
fiModelDescriptor
ProductSize
fiProductClassDesc
state
ProductGroup
ProductGroupDesc
Drive_System
Enclosure
Forks
Pad_Type
Ride_Control
Stick
Transmission
Turbocharged
Blade_Extension
Blade_Width
Enclosure_Type
Engine_Horsepower
Hydraulics
Pushblock
Ripper
Scarifier
Tip_Control
Tire_Size
Coupler
Coupler_System
Grouser_Tracks
Hydraulics_Flow
Track_Type
Undercarriage_Pad_Width
Stick_Length
Thumb
Pattern_Changer
Grouser_Type
Backhoe_Mounting
Blade_Type
Travel_Controls
Differential_Type
Steering_Controls

In [152]: *# This will turn all the string values into categorical values*

```
for label, content in df_tmp.items():
    if pd.api.types.is_string_dtype(content):
        df_tmp[label]=content.astype("category").cat.as_ordered()
```

```
In [153]: df_tmp.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 412698 entries, 205615 to 409203
Data columns (total 57 columns):
#   Column                                Non-Null Count  Dtype
---  ---                                -
0   SalesID                             412698 non-null  int64
1   SalePrice                           412698 non-null  float64
2   MachineID                           412698 non-null  int64
3   ModelID                             412698 non-null  int64
4   datasource                          412698 non-null  int64
5   auctioneerID                       392562 non-null  float64
6   YearMade                           412698 non-null  int64
7   MachineHoursCurrentMeter           147504 non-null  float64
8   UsageBand                          73670 non-null   category
9   fiModelDesc                        412698 non-null  category
10  fiBaseModel                        412698 non-null  category
11  fiSecondaryDesc                    271971 non-null  category
12  fiModelSeries                      58667 non-null   category
13  fiModelDescriptor                  74816 non-null   category
14  ProductSize                        196093 non-null  category
15  fiProductClassDesc                412698 non-null  category
16  state                              412698 non-null  category
17  ProductGroup                      412698 non-null  category
18  ProductGroupDesc                  412698 non-null  category
19  Drive_System                      107087 non-null  category
20  Enclosure                         412364 non-null  category
21  Forks                             197715 non-null  category
22  Pad_Type                          81096 non-null   category
23  Ride_Control                      152728 non-null  category
24  Stick                             81096 non-null   category
25  Transmission                      188007 non-null  category
26  Turbocharged                      81096 non-null   category
27  Blade_Extension                    25983 non-null   category
28  Blade_Width                       25983 non-null   category
29  Enclosure_Type                    25983 non-null   category
30  Engine_Horsepower                 25983 non-null   category
31  Hydraulics                        330133 non-null  category
32  Pushblock                         25983 non-null   category
33  Ripper                            106945 non-null  category
34  Scarifier                         25994 non-null   category
35  Tip_Control                       25983 non-null   category
36  Tire_Size                         97638 non-null   category
37  Coupler                           220679 non-null  category
38  Coupler_System                    44974 non-null   category
39  Grouser_Tracks                     44875 non-null   category
40  Hydraulics_Flow                   44875 non-null   category
41  Track_Type                        102193 non-null  category
42  Undercarriage_Pad_Width           102916 non-null  category
43  Stick_Length                      102261 non-null  category
44  Thumb                             102332 non-null  category
45  Pattern_Changer                   102261 non-null  category
46  Grouser_Type                      102193 non-null  category
47  Backhoe_Mounting                  80712 non-null   category
48  Blade_Type                        81875 non-null   category
49  Travel_Controls                   81877 non-null   category
```

```

50 Differential_Type      71564 non-null  category
51 Steering_Controls     71522 non-null  category
52 saleYear              412698 non-null int64
53 saleMonth             412698 non-null int64
54 saleDay               412698 non-null int64
55 saleDayOfWeek         412698 non-null int64
56 saleDayOfYear         412698 non-null int64
dtypes: category(44), float64(3), int64(10)
memory usage: 63.3 MB

```

```
In [154]: df_tmp.state.cat.categories
```

```

Out[154]: Index(['Alabama', 'Alaska', 'Arizona', 'Arkansas', 'California', 'Colorado',
                'Connecticut', 'Delaware', 'Florida', 'Georgia', 'Hawaii', 'Idaho',
                'Illinois', 'Indiana', 'Iowa', 'Kansas', 'Kentucky', 'Louisiana',
                'Maine', 'Maryland', 'Massachusetts', 'Michigan', 'Minnesota',
                'Mississippi', 'Missouri', 'Montana', 'Nebraska', 'Nevada',
                'New Hampshire', 'New Jersey', 'New Mexico', 'New York',
                'North Carolina', 'North Dakota', 'Ohio', 'Oklahoma', 'Oregon',
                'Pennsylvania', 'Puerto Rico', 'Rhode Island', 'South Carolina',
                'South Dakota', 'Tennessee', 'Texas', 'Unspecified', 'Utah', 'Vermont',
                'Virginia', 'Washington', 'Washington DC', 'West Virginia', 'Wisconsin',
                'Wyoming'],
                dtype='object')

```

```
In [155]: df_tmp.state.cat.codes
```

```

Out[155]: 205615      43
          274835      8
          141296      8
          212552      8
          62755      8
          ..
          410879      4
          412476      4
          411927      4
          407124      4
          409203      4
Length: 412698, dtype: int8

```

Save preprocessed data

```

In [156]: # Export current tmp dataframe
df_tmp.to_csv("data/train_tmp.csv", index=False)

```

```

In [157]: # Import preprocessed data
df_tmp=pd.read_csv("data/train_tmp.csv",
                  low_memory=False)

```

In [158]: df_tmp.head().T

Out[158]:

	0	1	2	3	4
SalesID	1646770	1821514	1505138	1671174	1329056
SalePrice	9500	14000	50000	16000	22000
MachineID	1126363	1194089	1473654	1327630	1336053
ModelID	8434	10150	4139	8591	4089
datasource	132	132	132	132	132
auctioneerID	18	99	99	99	99
YearMade	1974	1980	1978	1980	1984
MachineHoursCurrentMeter	NaN	NaN	NaN	NaN	NaN
UsageBand	NaN	NaN	NaN	NaN	NaN
fiModelDesc	TD20	A66	D7G	A62	D3B
fiBaseModel	TD20	A66	D7	A62	D3
fiSecondaryDesc	NaN	NaN	G	NaN	B
fiModelSeries	NaN	NaN	NaN	NaN	NaN
fiModelDescriptor	NaN	NaN	NaN	NaN	NaN
ProductSize	Medium	NaN	Large	NaN	NaN
fiProductClassDesc	Track Type Tractor, Dozer - 105.0 to 130.0 Hor...	Wheel Loader - 120.0 to 135.0 Horsepower	Track Type Tractor, Dozer - 190.0 to 260.0 Hor...	Wheel Loader - Unidentified	Track Type Tractor, Dozer - 20.0 to 75.0 Horse...
state	Texas	Florida	Florida	Florida	Florida
ProductGroup	TTT	WL	TTT	WL	TTT
ProductGroupDesc	Track Type Tractors	Wheel Loader	Track Type Tractors	Wheel Loader	Track Type Tractors
Drive_System	NaN	NaN	NaN	NaN	NaN
Enclosure	OROPS	OROPS	OROPS	EROPS	OROPS
Forks	NaN	None or Unspecified	NaN	None or Unspecified	NaN
Pad_Type	NaN	NaN	NaN	NaN	NaN
Ride_Control	NaN	None or Unspecified	NaN	None or Unspecified	NaN
Stick	NaN	NaN	NaN	NaN	NaN
Transmission	Direct Drive	NaN	Standard	NaN	Standard
Turbocharged	NaN	NaN	NaN	NaN	NaN
Blade_Extension	NaN	NaN	NaN	NaN	NaN
Blade_Width	NaN	NaN	NaN	NaN	NaN
Enclosure_Type	NaN	NaN	NaN	NaN	NaN

	0	1	2	3	4
Engine_Horsepower	NaN	NaN	NaN	NaN	NaN
Hydraulics	2 Valve	2 Valve	2 Valve	2 Valve	2 Valve
Pushblock	NaN	NaN	NaN	NaN	NaN
Ripper	None or Unspecified	NaN	None or Unspecified	NaN	None or Unspecified
Scarifier	NaN	NaN	NaN	NaN	NaN
Tip_Control	NaN	NaN	NaN	NaN	NaN
Tire_Size	NaN	None or Unspecified	NaN	None or Unspecified	NaN
Coupler	NaN	None or Unspecified	NaN	None or Unspecified	NaN
Coupler_System	NaN	NaN	NaN	NaN	NaN
Grouser_Tracks	NaN	NaN	NaN	NaN	NaN
Hydraulics_Flow	NaN	NaN	NaN	NaN	NaN
Track_Type	NaN	NaN	NaN	NaN	NaN
Undercarriage_Pad_Width	NaN	NaN	NaN	NaN	NaN
Stick_Length	NaN	NaN	NaN	NaN	NaN
Thumb	NaN	NaN	NaN	NaN	NaN
Pattern_Changer	NaN	NaN	NaN	NaN	NaN
Grouser_Type	NaN	NaN	NaN	NaN	NaN
Backhoe_Mounting	None or Unspecified	NaN	None or Unspecified	NaN	None or Unspecified
Blade_Type	Straight	NaN	Straight	NaN	PAT
Travel_Controls	None or Unspecified	NaN	None or Unspecified	NaN	Lever
Differential_Type	NaN	Standard	NaN	Standard	NaN
Steering_Controls	NaN	Conventional	NaN	Conventional	NaN
saleYear	1989	1989	1989	1989	1989
saleMonth	1	1	1	1	1
saleDay	17	31	31	31	31
saleDayOfWeek	1	1	1	1	1
saleDayOfYear	17	31	31	31	31

```
In [159]: df_tmp.isnull().sum()
```

```
Out[159]: SalesID                0
SalePrice                0
MachineID                0
ModelID                  0
datasource                0
auctioneerID             20136
YearMade                  0
MachineHoursCurrentMeter 265194
UsageBand                 339028
fiModelDesc               0
fiBaseModel               0
fiSecondaryDesc           140727
fiModelSeries             354031
fiModelDescriptor         337882
ProductSize               216605
fiProductClassDesc        0
state                     0
ProductGroup              0
ProductGroupDesc          0
Drive_System              305611
Enclosure                 334
Forks                     214983
Pad_Type                  331602
Ride_Control              259970
Stick                     331602
Transmission              224691
Turbocharged              331602
Blade_Extension           386715
Blade_Width               386715
Enclosure_Type            386715
Engine_Horsepower         386715
Hydraulics                82565
Pushblock                 386715
Ripper                    305753
Scarifier                 386704
Tip_Control               386715
Tire_Size                 315060
Coupler                   192019
Coupler_System            367724
Grouser_Tracks            367823
Hydraulics_Flow           367823
Track_Type                310505
Undercarriage_Pad_Width   309782
Stick_Length              310437
Thumb                     310366
Pattern_Changer           310437
Grouser_Type              310505
Backhoe_Mounting          331986
Blade_Type                330823
Travel_Controls           330821
Differential_Type         341134
Steering_Controls         341176
saleYear                  0
saleMonth                 0
saleDay                   0
```



```
saleDayOfWeek      0
saleDayOfYear       0
dtype: int64
```

Fill the missing values

First fill the numeric missing values

```
In [160]: # First check the numerical columns
for label,content in df_tmp.items():
    if pd.api.types.is_numeric_dtype(content):
        print(label)
```

```
SalesID
SalePrice
MachineID
ModelID
datasource
auctioneerID
YearMade
MachineHoursCurrentMeter
saleYear
saleMonth
saleDay
saleDayOfWeek
saleDayOfYear
```

```
In [161]: # Check for which numeric columns have null values
for label, content in df_tmp.items():
    if pd.api.types.is_numeric_dtype(content):
        if pd.isnull(content).sum():
            print(label)
```

```
auctioneerID
MachineHoursCurrentMeter
```

```
In [162]: # Fill numeric rows with median
for label, content in df_tmp.items():
    if pd.api.types.is_numeric_dtype(content):
        if pd.isnull(content).sum():
            # Add a binary column which tells us if the data was missing or not
            df_tmp[label+"_is_missing"]=pd.isnull(content)
            # Fill missing numeric values with median
            df_tmp[label]=content.fillna(content.median())
```

```
In [164]: # Check if there is null numeric value
for label, content in df_tmp.items():
    if pd.api.types.is_numeric_dtype(content):
        if pd.isnull(content).sum():
            print(label)
```

```
In [165]: # Checkk to see how many examples were missing  
df_tmp.auctioneerID_is_missing.value_counts()
```

```
Out[165]: False      392562  
          True       20136  
          Name: auctioneerID_is_missing, dtype: int64
```

```
In [166]: df_tmp.isna().sum()
```

```
Out[166]: SalesID                0
SalePrice                0
MachineID                0
ModelID                 0
datasource              0
auctioneerID            0
YearMade                0
MachineHoursCurrentMeter 0
UsageBand              339028
fiModelDesc             0
fiBaseModel             0
fiSecondaryDesc         140727
fiModelSeries           354031
fiModelDescriptor       337882
ProductSize            216605
fiProductClassDesc      0
state                  0
ProductGroup            0
ProductGroupDesc        0
Drive_System            305611
Enclosure               334
Forks                  214983
Pad_Type               331602
Ride_Control           259970
Stick                  331602
Transmission           224691
Turbocharged           331602
Blade_Extension         386715
Blade_Width            386715
Enclosure_Type         386715
Engine_Horsepower       386715
Hydraulics              82565
Pushblock              386715
Ripper                 305753
Scarifier              386704
Tip_Control            386715
Tire_Size              315060
Coupler                192019
Coupler_System         367724
Grouser_Tracks         367823
Hydraulics_Flow        367823
Track_Type             310505
Undercarriage_Pad_Width 309782
Stick_Length           310437
Thumb                  310366
Pattern_Changer        310437
Grouser_Type           310505
Backhoe_Mounting       331986
Blade_Type             330823
Travel_Controls        330821
Differential_Type      341134
Steering_Controls      341176
saleYear               0
saleMonth              0
saleDay                0
```

```
saleDayOfWeek          0
saleDayOfYear           0
auctioneerID_is_missing 0
MachineHoursCurrentMeter_is_missing 0
dtype: int64
```

Filling and turning categorical values with numerical values

```
In [167]: # Find the categorical column with null values
for label, content in df_tmp.items():
    if not pd.api.types.is_numeric_dtype(content):
        if pd.isnull(content).sum():
            print(label)
```

```
UsageBand
fiSecondaryDesc
fiModelSeries
fiModelDescriptor
ProductSize
Drive_System
Enclosure
Forks
Pad_Type
Ride_Control
Stick
Transmission
Turbocharged
Blade_Extension
Blade_Width
Enclosure_Type
Engine_Horsepower
Hydraulics
Pushblock
Ripper
Scarifier
Tip_Control
Tire_Size
Coupler
Coupler_System
Grouser_Tracks
Hydraulics_Flow
Track_Type
Undercarriage_Pad_Width
Stick_Length
Thumb
Pattern_Changer
Grouser_Type
Backhoe_Mounting
Blade_Type
Travel_Controls
Differential_Type
Steering_Controls
```

```
In [168]: pd.Categorical(df_tmp["UsageBand"]).codes+1
```

```
Out[168]: array([0, 0, 0, ..., 0, 0, 0], dtype=int8)
```

```
In [169]: # Turn categorical values into number and fill missing values
for label, content in df_tmp.items():
    if not pd.api.types.is_numeric_dtype(content):
        # Add binary columns to indicate whether sample had missing values or not
        df_tmp[label+"_is_missing"]=pd.isnull(content)
        # Turn categories into numbers and add +1
        df_tmp[label]=pd.Categorical(content).codes+1
```

```
In [170]: df_tmp.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 412698 entries, 0 to 412697
Columns: 103 entries, SalesID to Steering_Controls_is_missing
dtypes: bool(46), float64(3), int16(4), int64(10), int8(40)
memory usage: 77.9 MB
```

```
In [171]: df_tmp.head().T
```

```
Out[171]:
```

	0	1	2	3	4
SalesID	1646770	1821514	1505138	1671174	1329056
SalePrice	9500	14000	50000	16000	22000
MachineID	1126363	1194089	1473654	1327630	1336053
ModelID	8434	10150	4139	8591	4089
datasource	132	132	132	132	132
...
Backhoe_Mounting_is_missing	False	True	False	True	False
Blade_Type_is_missing	False	True	False	True	False
Travel_Controls_is_missing	False	True	False	True	False
Differential_Type_is_missing	True	False	True	False	True
Steering_Controls_is_missing	True	False	True	False	True

103 rows × 5 columns

```
In [172]: df_tmp.isna().sum()
```

```
Out[172]: SalesID          0
          SalePrice        0
          MachineID        0
          ModelID          0
          datasource        0
          ..
          Backhoe_Mounting_is_missing  0
          Blade_Type_is_missing        0
          Travel_Controls_is_missing   0
          Differential_Type_is_missing  0
          Steering_Controls_is_missing  0
          Length: 103, dtype: int64
```

Now let's build the Machine Learning model

```
In [173]: len(df_tmp)
```


```
Out[173]: 412698
```

```
In [174]: df_tmp.head()
```

```
Out[174]:
```

	SalesID	SalePrice	MachineID	ModelID	datasource	auctioneerID	YearMade	MachineHoursCur
0	1646770	9500.0	1126363	8434	132	18.0	1974	
1	1821514	14000.0	1194089	10150	132	99.0	1980	
2	1505138	50000.0	1473654	4139	132	99.0	1978	
3	1671174	16000.0	1327630	8591	132	99.0	1980	
4	1329056	22000.0	1336053	4089	132	99.0	1984	

5 rows × 103 columns



```
In [175]: %%time
          # Instantiate model
          model = RandomForestRegressor(n_jobs=-1,
                                       random_state=42)

          # Fit the model
          model.fit(df_tmp.drop("SalePrice", axis=1), df_tmp["SalePrice"])
```

Wall time: 7min 44s

```
Out[175]: RandomForestRegressor(n_jobs=-1, random_state=42)
```

```
In [176]: # Score the model
          model.score(df_tmp.drop("SalePrice", axis=1), df_tmp["SalePrice"])
```

```
Out[176]: 0.9875468079970562
```

Split the data in train and valid sets (Plzz see the data)

```
In [177]: df_tmp.saleYear
```

```
Out[177]: 0          1989
          1          1989
          2          1989
          3          1989
          4          1989
          ...
         412693      2012
         412694      2012
         412695      2012
         412696      2012
         412697      2012
          Name: saleYear, Length: 412698, dtype: int64
```

```
In [178]: df_tmp.saleYear.value_counts()
```

```
Out[178]: 2009      43849
          2008      39767
          2011      35197
          2010      33390
          2007      32208
          2006      21685
          2005      20463
          2004      19879
          2001      17594
          2000      17415
          2002      17246
          2003      15254
          1998      13046
          1999      12793
          2012      11573
          1997       9785
          1996       8829
          1995       8530
          1994       7929
          1993       6303
          1992       5519
          1991       5109
          1989       4806
          1990       4529
          Name: saleYear, dtype: int64
```

```
In [179]: df_valid = df_tmp[df_tmp.saleYear==2012]
          df_train = df_tmp[df_tmp.saleYear!=2012]
          len(df_valid), len(df_train)
```

```
Out[179]: (11573, 401125)
```

```
In [180]: # Split the train and valid data in X and y
X_train, y_train = df_train.drop("SalePrice", axis=1), df_train["SalePrice"]
X_valid, y_valid = df_valid.drop("SalePrice", axis=1), df_valid["SalePrice"]
X_train.shape, y_train.shape, X_valid.shape, y_valid.shape
```

```
Out[180]: ((401125, 102), (401125,), (11573, 102), (11573,))
```

```
In [181]: y_train
```

```
Out[181]: 0          9500.0
1         14000.0
2         50000.0
3         16000.0
4         22000.0
...
401120     29000.0
401121     11000.0
401122     11000.0
401123     18000.0
401124     13500.0
Name: SalePrice, Length: 401125, dtype: float64
```

Build an Evaluation Metrics

```
In [182]: # Create a evaluation metrics i.e RMSLE
from sklearn.metrics import mean_squared_log_error, mean_absolute_error, r2_score

def rmsle(y_test, y_preds):
    """
    Calculate root mean squared log error between prediction and true labels.
    """
    return np.sqrt(mean_squared_log_error(y_test, y_preds))

# Create function to evaluate model on a few different levels.
def show_score(model):
    train_preds=model.predict(X_train)
    valid_preds=model.predict(X_valid)
    scores={"Training MAE":mean_absolute_error(y_train,train_preds),
            "Valid MAE":mean_absolute_error(y_valid,valid_preds),
            "Training RMSLE":rmsle(y_train,train_preds),
            "Valid RMSLE":rmsle(y_valid,valid_preds),
            "Training R^2":r2_score(y_train,train_preds),
            "Valid R^2":r2_score(y_valid,valid_preds)}
    return scores
```

Testing our model on Subset (to tune the hyperparameters)


```
In [183]: # This takes far too long .... for experimenting

# %%time
# model=RandomForestRegressor(n_jobs=-1,
#                               random_state=42)
# model.fit(X_train,y_train)
```

```
In [184]: len(X_train)
```

```
Out[184]: 401125
```

```
In [185]: # Change max samples value
from sklearn.ensemble import RandomForestRegressor
model=RandomForestRegressor(n_jobs=-1,
                             random_state=42,
                             max_samples=10000)
```

```
In [186]: %%time
# Cutting down on the max number of samples each estimator can see improve training
model.fit(X_train, y_train)
```

Wall time: 25.4 s

```
Out[186]: RandomForestRegressor(max_samples=10000, n_jobs=-1, random_state=42)
```

```
In [187]: show_score(model)
```

```
Out[187]: {'Training MAE': 5561.2988092240585,
'Valid MAE': 7177.26365505919,
'Training RMSLE': 0.257745378256977,
'Valid RMSLE': 0.29362638671089003,
'Training R^2': 0.8606658995199189,
'Valid R^2': 0.8320374995090507}
```

Hyperparameter tuning with RandomizedSearchCV

```

In [188]: %%time
from sklearn.model_selection import RandomizedSearchCV

# Different RandomForestRegressor hyperparameters
rf_grid={"n_estimators":np.arange(10,100,10),
        "max_depth":[None,3,5,10],
        "min_samples_split":np.arange(2,20,2),
        "min_samples_leaf":np.arange(1,20,2),
        "max_features":[0.5,1,"sqrt","auto"],
        "max_samples":[10000]}

# Instantiate RandomizedSearchCV model
rs_model = RandomizedSearchCV(RandomForestRegressor(n_jobs=-1,
                                                    random_state=42),
                              param_distributions=rf_grid,
                              n_iter=2,
                              cv=5,
                              verbose=True)

# Fit the model
rs_model.fit(X_train,y_train)

```

Fitting 5 folds for each of 2 candidates, totalling 10 fits
Wall time: 3min 25s

```

Out[188]: RandomizedSearchCV(cv=5,
                             estimator=RandomForestRegressor(n_jobs=-1, random_state=42),
                             n_iter=2,
                             param_distributions={'max_depth': [None, 3, 5, 10],
                                                  'max_features': [0.5, 1, 'sqrt',
                                                                'auto'],
                                                  'max_samples': [10000],
                                                  'min_samples_leaf': array([ 1,  3,  5,
7,  9, 11, 13, 15, 17, 19]),
                                                  'min_samples_split': array([ 2,  4,  6,
8, 10, 12, 14, 16, 18]),
                                                  'n_estimators': array([10, 20, 30, 40,
50, 60, 70, 80, 90])},
                             verbose=True)

```

```

In [189]: # Find the best model hyperparameters
rs_model.best_params_

```

```

Out[189]: {'n_estimators': 40,
           'min_samples_split': 12,
           'min_samples_leaf': 9,
           'max_samples': 10000,
           'max_features': 'auto',
           'max_depth': None}

```

```
In [190]: # Evaluate the RandomizedSearchCV model
show_score(rs_model)
```

```
Out[190]: {'Training MAE': 6086.99131166347,
'Valid MAE': 7674.949511960607,
'Training RMSLE': 0.27648998624118665,
'Valid RMSLE': 0.3072434602008128,
'Training R^2': 0.8312379190695804,
'Valid R^2': 0.8002384558444482}
```

Train a model with the best hyperparameters

Note: These were found after 100 iterations of RandomizedSearchCV

```
In [191]: %%time

# Most ideal hyperparameter
ideal_model = RandomForestRegressor(n_estimators=40,
                                   min_samples_leaf=1,
                                   min_samples_split=14,
                                   max_features=0.5,
                                   n_jobs=-1,
                                   max_samples=None,
                                   random_state=42)

# Fit the ideal model
ideal_model.fit(X_train,y_train)
```

Wall time: 1min 34s

```
Out[191]: RandomForestRegressor(max_features=0.5, min_samples_split=14, n_estimators=40,
                                n_jobs=-1, random_state=42)
```

```
In [192]: # Scores for ideal model (trained on all the data)
show_score(ideal_model)
```

```
Out[192]: {'Training MAE': 2953.8161137163484,
'Valid MAE': 5951.247761444453,
'Training RMSLE': 0.14469006962371858,
'Valid RMSLE': 0.24524163989538328,
'Training R^2': 0.9588145522577225,
'Valid R^2': 0.8818019502450094}
```

```
In [193]: # Scores for rs_model (only trained on 10,000 examples)
show_score(rs_model)
```

```
Out[193]: {'Training MAE': 6086.99131166347,
'Valid MAE': 7674.949511960607,
'Training RMSLE': 0.27648998624118665,
'Valid RMSLE': 0.3072434602008128,
'Training R^2': 0.8312379190695804,
'Valid R^2': 0.8002384558444482}
```

Test our model using test data

```
In [203]: # Import the test data
df_test=pd.read_csv("data/Test.csv",
                    low_memory=False,
                    parse_dates=["saledate"])
df_test.head().T
```

```
Out[203]:
```

	0	1	2	3	4
SalesID	1227829	1227844	1227847	1227848	1227863
MachineID	1006309	1022817	1031560	56204	1053887
ModelID	3168	7271	22805	1269	22312
datasource	121	121	121	121	121
auctioneerID	3	3	3	3	3
YearMade	1999	1000	2004	2006	2005
MachineHoursCurrentMeter	3688	28555	6038	8940	2286
UsageBand	Low	High	Medium	High	Low
saledate	2012-05-03 00:00:00	2012-05-10 00:00:00	2012-05-10 00:00:00	2012-05-10 00:00:00	2012-05-10 00:00:00
fiModelDesc	580G	936	EC210BLC	330CL	650K
fiBaseModel	580	936	EC210	330	650
fiSecondaryDesc	G	NaN	B	C	K
fiModelSeries	NaN	NaN	NaN	NaN	NaN
fiModelDescriptor	NaN	NaN	LC	L	NaN
ProductSize	NaN	Medium	Large / Medium	Large / Medium	NaN
fiProductClassDesc	Backhoe Loader - 14.0 to 15.0 Ft Standard Digg...	Wheel Loader - 135.0 to 150.0 Horsepower	Hydraulic Excavator, Track - 21.0 to 24.0 Metr...	Hydraulic Excavator, Track - 33.0 to 40.0 Metr...	Track Type Tractor, Dozer - 20.0 to 75.0 Horse...
state	Wyoming	Virginia	New Jersey	New Jersey	Florida
ProductGroup	BL	WL	TEX	TEX	TTT
ProductGroupDesc	Backhoe Loaders	Wheel Loader	Track Excavators	Track Excavators	Track Type Tractors
Drive_System	Two Wheel Drive	NaN	NaN	NaN	NaN
Enclosure	OROPS	EROPS	EROPS w AC	EROPS w AC	OROPS
Forks	Yes	Yes	NaN	NaN	NaN
Pad_Type	None or Unspecified	NaN	NaN	NaN	NaN
Ride_Control	No	None or Unspecified	NaN	NaN	NaN
Stick	Standard	NaN	NaN	NaN	NaN
Transmission	Standard	NaN	NaN	NaN	Hydrostatic

	0	1	2	3	4
Turbocharged	None or Unspecified	NaN	NaN	NaN	NaN
Blade_Extension	NaN	NaN	NaN	NaN	NaN
Blade_Width	NaN	NaN	NaN	NaN	NaN
Enclosure_Type	NaN	NaN	NaN	NaN	NaN
Engine_Horsepower	NaN	NaN	NaN	NaN	NaN
Hydraulics	NaN	2 Valve	Auxiliary	Standard	2 Valve
Pushblock	NaN	NaN	NaN	NaN	NaN
Ripper	NaN	NaN	NaN	NaN	None or Unspecified
Scarifier	NaN	NaN	NaN	NaN	NaN
Tip_Control	NaN	NaN	NaN	NaN	NaN
Tire_Size	NaN	20.5	NaN	NaN	NaN
Coupler	NaN	None or Unspecified	None or Unspecified	None or Unspecified	NaN
Coupler_System	NaN	NaN	NaN	NaN	NaN
Grouser_Tracks	NaN	NaN	NaN	NaN	NaN
Hydraulics_Flow	NaN	NaN	NaN	NaN	NaN
Track_Type	NaN	NaN	Steel	Steel	NaN
Undercarriage_Pad_Width	NaN	NaN	None or Unspecified	None or Unspecified	NaN
Stick_Length	NaN	NaN	9' 6"	None or Unspecified	NaN
Thumb	NaN	NaN	Manual	Manual	NaN
Pattern_Changer	NaN	NaN	None or Unspecified	Yes	NaN
Grouser_Type	NaN	NaN	Double	Triple	NaN
Backhoe_Mounting	NaN	NaN	NaN	NaN	None or Unspecified
Blade_Type	NaN	NaN	NaN	NaN	PAT
Travel_Controls	NaN	NaN	NaN	NaN	None or Unspecified
Differential_Type	NaN	Standard	NaN	NaN	NaN
Steering_Controls	NaN	Conventional	NaN	NaN	NaN

```
In [204]: df_test.saledate[:1000]
```

```
Out[204]: 0      2012-05-03
          1      2012-05-10
          2      2012-05-10
          3      2012-05-10
          4      2012-05-10
          ...
          995    2012-06-12
          996    2012-06-12
          997    2012-05-01
          998    2012-05-01
          999    2012-06-05
          Name: saledate, Length: 1000, dtype: datetime64[ns]
```

```
In [205]: df_test.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 12457 entries, 0 to 12456
Data columns (total 52 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   SalesID                               12457 non-null  int64
1   MachineID                             12457 non-null  int64
2   ModelID                               12457 non-null  int64
3   datasource                            12457 non-null  int64
4   auctioneerID                          12457 non-null  int64
5   YearMade                              12457 non-null  int64
6   MachineHoursCurrentMeter              2129 non-null   float64
7   UsageBand                             1834 non-null   object
8   saledate                              12457 non-null  datetime64[ns]
9   fiModelDesc                           12457 non-null  object
10  fiBaseModel                           12457 non-null  object
11  fiSecondaryDesc                        8482 non-null   object
12  fiModelSeries                          2006 non-null   object
13  fiModelDescriptor                      3024 non-null   object
14  ProductSize                           6048 non-null   object
15  fiProductClassDesc                    12457 non-null  object
16  state                                 12457 non-null  object
17  ProductGroup                          12457 non-null  object
18  ProductGroupDesc                      12457 non-null  object
19  Drive_System                          2759 non-null   object
20  Enclosure                             12455 non-null  object
21  Forks                                 6308 non-null   object
22  Pad_Type                              2108 non-null   object
23  Ride_Control                          4241 non-null   object
24  Stick                                 2108 non-null   object
25  Transmission                          4818 non-null   object
26  Turbocharged                          2108 non-null   object
27  Blade_Extension                       651 non-null    object
28  Blade_Width                           651 non-null    object
29  Enclosure_Type                        651 non-null    object
30  Engine_Horsepower                     651 non-null    object
31  Hydraulics                            10315 non-null  object
32  Pushblock                             651 non-null    object
33  Ripper                                2704 non-null   object
34  Scarifier                             651 non-null    object
35  Tip_Control                           651 non-null    object
36  Tire_Size                             2778 non-null   object
37  Coupler                               7601 non-null   object
38  Coupler_System                        2066 non-null   object
39  Grouser_Tracks                        2066 non-null   object
40  Hydraulics_Flow                       2066 non-null   object
41  Track_Type                             3394 non-null   object
42  Undercarriage_Pad_Width                3398 non-null   object
43  Stick_Length                           3394 non-null   object
44  Thumb                                 3395 non-null   object
45  Pattern_Changer                        3394 non-null   object
46  Grouser_Type                           3394 non-null   object
47  Backhoe_Mounting                       2051 non-null   object
48  Blade_Type                             2058 non-null   object
49  Travel_Controls                       2058 non-null   object
```



```
50 Differential_Type          2129 non-null  object
51 Steering_Controls          2129 non-null  object
dtypes: datetime64[ns](1), float64(1), int64(6), object(44)
memory usage: 4.9+ MB
```

```
In [206]: df_test.isna().sum()
```

```
Out[206]: SalesID                0
MachineID                0
ModelID                  0
datasource                0
auctioneerID              0
YearMade                  0
MachineHoursCurrentMeter 10328
UsageBand                 10623
saledate                  0
fiModelDesc               0
fiBaseModel               0
fiSecondaryDesc           3975
fiModelSeries             10451
fiModelDescriptor         9433
ProductSize               6409
fiProductClassDesc        0
state                     0
ProductGroup              0
ProductGroupDesc          0
Drive_System              9698
Enclosure                 2
Forks                     6149
Pad_Type                  10349
Ride_Control              8216
Stick                     10349
Transmission              7639
Turbocharged              10349
Blade_Extension           11806
Blade_Width               11806
Enclosure_Type            11806
Engine_Horsepower         11806
Hydraulics                2142
Pushblock                 11806
Ripper                    9753
Scarifier                 11806
Tip_Control               11806
Tire_Size                 9679
Coupler                   4856
Coupler_System            10391
Grouser_Tracks            10391
Hydraulics_Flow           10391
Track_Type                9063
Undercarriage_Pad_Width   9059
Stick_Length              9063
Thumb                     9062
Pattern_Changer           9063
Grouser_Type              9063
Backhoe_Mounting          10406
Blade_Type                10399
Travel_Controls           10399
Differential_Type         10328
Steering_Controls         10328
dtype: int64
```

Preprocessing the data (getting the test dataset in the same format as our training dataset)

```
In [207]: def preprocess_data(df):  
    """  
    Performs some transformation on df and returns transformed df.  
    """  
    df["saleYear"] = df.saledate.dt.year  
    df["saleMonth"] = df.saledate.dt.month  
    df["saleDay"] = df.saledate.dt.day  
    df["saleDayOfWeek"] = df.saledate.dt.dayofweek  
    df["saleDayOfYear"] = df.saledate.dt.dayofyear  
  
    df.drop(labels="saledate", axis=1, inplace=True)  
  
    # Fill the numeric rows with median  
    for label, content in df.items():  
        if pd.api.types.is_numeric_dtype(content):  
            if pd.isnull(content).sum():  
                # Add a binary column which tells us if the data was missing or not  
                df[label + "_is_missing"] = pd.isnull(content)  
                # Fill missing numeric values with median  
                df[label] = content.fillna(content.median())  
  
    # Fill the categorical missing data and turn categories into numbers  
    if not pd.api.types.is_numeric_dtype(content):  
        df[label + "_is_missing"] = pd.isnull(content)  
        df[label] = pd.Categorical(content).codes + 1  
  
    return df
```

```
In [208]: # Preprocess the test data  
df_test = preprocess_data(df_test)  
df_test.head()
```

```
Out[208]:
```

	SalesID	MachineID	ModelID	datasource	auctioneerID	YearMade	MachineHoursCurrentMeter
0	1227829	1006309	3168	121	3	1999	3688.0
1	1227844	1022817	7271	121	3	1000	28555.0
2	1227847	1031560	22805	121	3	2004	6038.0
3	1227848	56204	1269	121	3	2006	8940.0
4	1227863	1053887	22312	121	3	2005	2286.0

5 rows × 101 columns



```
In [209]: # Make predicytions on update test data
test_preds=ideal_model.predict(df_test)
```

```
-----
ValueError                                Traceback (most recent call last)
<ipython-input-209-2cc34114db25> in <module>
      1 # Make predicytions on update test data
----> 2 test_preds=ideal_model.predict(df_test)

~\AppData\Roaming\Python\Python37\site-packages\sklearn\ensemble\_forest.py in
predict(self, X)
    782         check_is_fitted(self)
    783         # Check data
--> 784         X = self._validate_X_predict(X)
    785
    786         # Assign chunk of trees to jobs

~\AppData\Roaming\Python\Python37\site-packages\sklearn\ensemble\_forest.py in
_validate_X_predict(self, X)
    420         check_is_fitted(self)
    421
--> 422         return self.estimators_[0]._validate_X_predict(X, check_input=Tr
rue)
    423
    424         @property

~\AppData\Roaming\Python\Python37\site-packages\sklearn\tree\_classes.py in _va
lidate_X_predict(self, X, check_input)
    406         if check_input:
    407             X = self._validate_data(X, dtype=DTYPE, accept_sparse="cs
r",
--> 408                                     reset=False)
    409             if issparse(X) and (X.indices.dtype != np.intc or
    410                               X.indptr.dtype != np.intc):

~\AppData\Roaming\Python\Python37\site-packages\sklearn\base.py in _validate_da
ta(self, X, y, reset, validate_separately, **check_params)
    435
    436         if check_params.get('ensure_2d', True):
--> 437             self._check_n_features(X, reset=reset)
    438
    439         return out

~\AppData\Roaming\Python\Python37\site-packages\sklearn\base.py in _check_n_fea
tures(self, X, reset)
    364         if n_features != self.n_features_in_:
    365             raise ValueError(
--> 366                 f"X has {n_features} features, but {self.__class__.__na
me__} "
    367                 f"is expecting {self.n_features_in_} features as input.
t.")
    368
```

ValueError: X has 101 features, but DecisionTreeRegressor is expecting 102 features as input.

```
In [210]: X_train.head()
```

```
Out[210]:
```

	SalesID	MachineID	ModelID	datasource	auctioneerID	YearMade	MachineHoursCurrentMeter
0	1646770	1126363	8434	132	18.0	1974	0.0
1	1821514	1194089	10150	132	99.0	1980	0.0
2	1505138	1473654	4139	132	99.0	1978	0.0
3	1671174	1327630	8591	132	99.0	1980	0.0
4	1329056	1336053	4089	132	99.0	1984	0.0

5 rows × 102 columns

```
In [211]: # we can find how the columns differ using sets
set(X_train.columns)-set(df_test.columns)
```

```
Out[211]: {'auctioneerID_is_missing'}
```

```
In [212]: # Manually adjust df_test to have auctioneerID_is_missing column
df_test["auctioneerID_is_missing"]=False
df_test.head()
```

```
Out[212]:
```

	SalesID	MachineID	ModelID	datasource	auctioneerID	YearMade	MachineHoursCurrentMeter
0	1227829	1006309	3168	121	3	1999	3688.0
1	1227844	1022817	7271	121	3	1000	28555.0
2	1227847	1031560	22805	121	3	2004	6038.0
3	1227848	56204	1269	121	3	2006	8940.0
4	1227863	1053887	22312	121	3	2005	2286.0

5 rows × 102 columns

```
In [213]: # Make predictions on the test data
test_preds=ideal_model.predict(df_test)
```

```
In [214]: test_preds
```

```
Out[214]: array([20614.36780887, 19897.80170658, 44852.21959446, ...,
14296.98620472, 22164.85757662, 31683.80063427])
```

```
In [215]: # Format predicitons into the same format kaggle is asking
df_preds=pd.DataFrame()
df_preds["SalesID"]=df_test["SalesID"]
df_preds["SalesPrice"]=test_preds
df_preds
```

```
Out[215]:
```

	SalesID	SalesPrice
0	1227829	20614.367809
1	1227844	19897.801707
2	1227847	44852.219594
3	1227848	68346.325323
4	1227863	39487.349708
...
12452	6643171	46466.092910
12453	6643173	17500.493352
12454	6643184	14296.986205
12455	6643186	22164.857577
12456	6643196	31683.800634

12457 rows × 2 columns

```
In [216]: # Export predictive data
df_preds.to_csv("data/test_predictions.csv", index=False)
```

```
In [ ]:
```